



# Neo4J Datalog Wrapper



# Plan



- ◇ Introduction
- ◇ General Requirements
- ◇ Project Structure
- ◇ Specific Team Requirements
- ◇ Tools
- ◇ Development Process
- ◇ Execution Plans
- ◇ Issues/Improvements
- ◇ Lessons/Future Outlook




# Introduction

Need for virtualized architecture



# Team Members



Julius Remigio

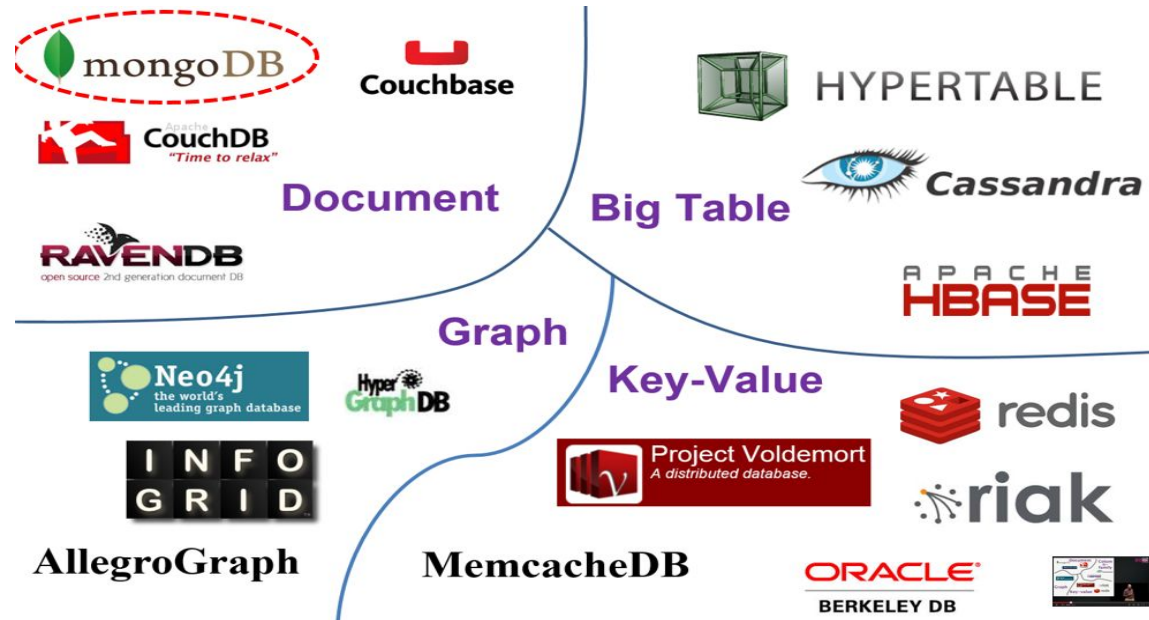
Ryan Riopelle

Michael Galarnyk





# DSE 203 – Importance of Data Integration



A cluster of decorative hexagonal icons in various shades of blue and teal. The icons include a lightbulb, a thumbs-up, a network node, a smartphone, a magnifying glass, a gear, and a speech bubble. The central hexagon is the largest and contains the number '2'.

2

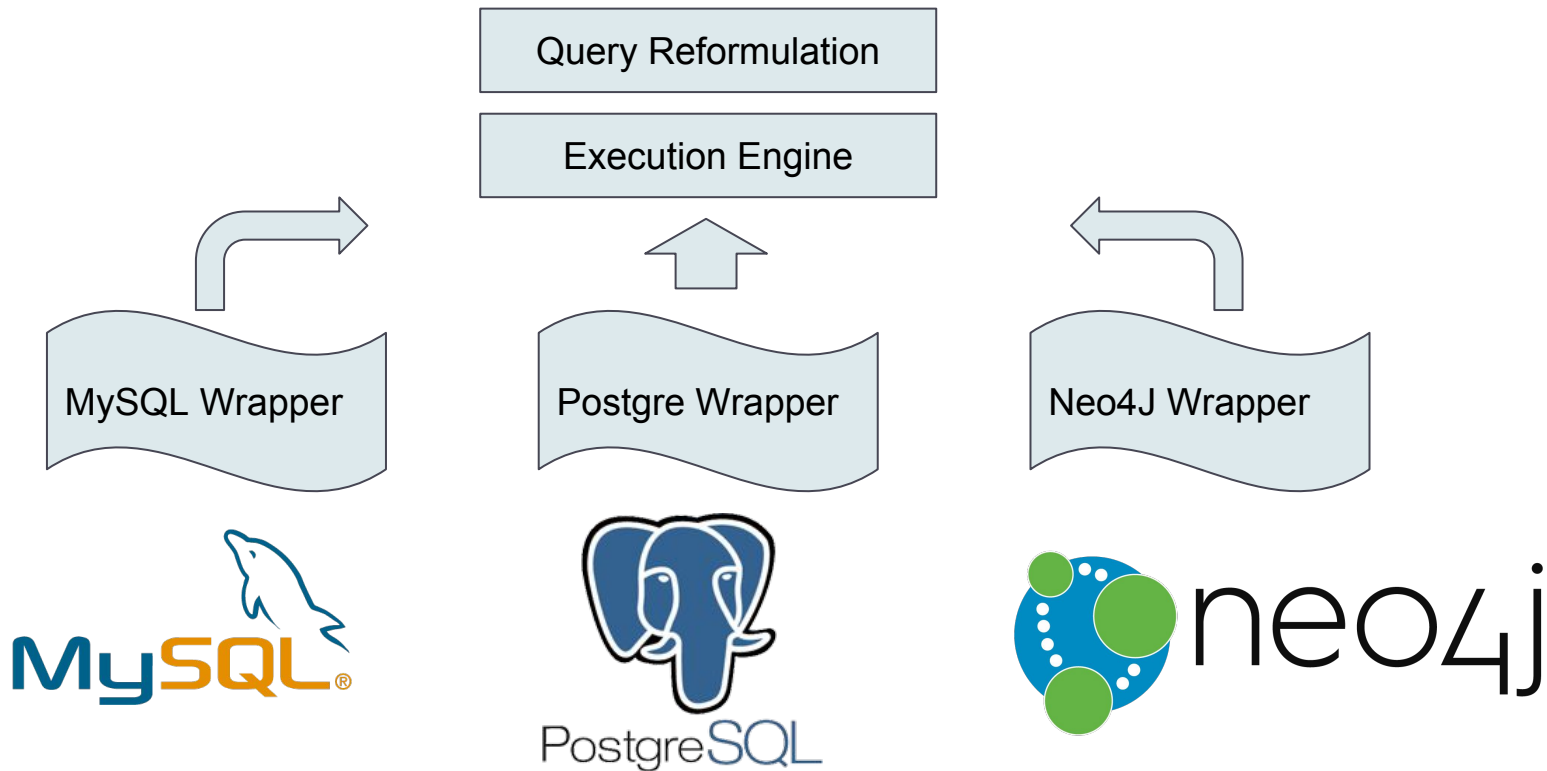
# Project Structure

Need for virtualized architecture

---




# Virtual Data Integration Architecture






# Query Languages




# Graph Databases

## Compared to Relational Databases

Optimized for aggregation



Optimized for connections



# Neo4J Wrapper Process Flow



q(organization) :-  
actor(id, \_, pname, \_),  
affiliation(id, organization, \_, \_),  
pname = 'Ariel Sharon'

Tables: ['affiliation', 'actor']  
Columns: ['organization', 'pname',  
'id']  
Required Projection: ['ization']  
Data is in Neo4J Schema A

Match (a: Actor {Name: 'Ariel  
Sharon'})-[aff:Affiliation]->(r)  
return r.Name as organization

		pname	id	ptype
0		Honorary Consul	64717	Group
1		Actor	64151	Individual
2		VIP	65561	Individual
3		Vips	65563	Group
4		Presidential Family	65216	Group
5		Retired	65338	Group
6		Infiltration Unit	64755	Group
7		Combatant	64411	Individual
8		Death Squad	64475	Group
9		Armed Professional	64226	Individual
10		Armed Force	64216	Group
11		Armed Gang	64218	Group
12		Armed Band	64214	Group



# Tools

Need for virtualized architecture

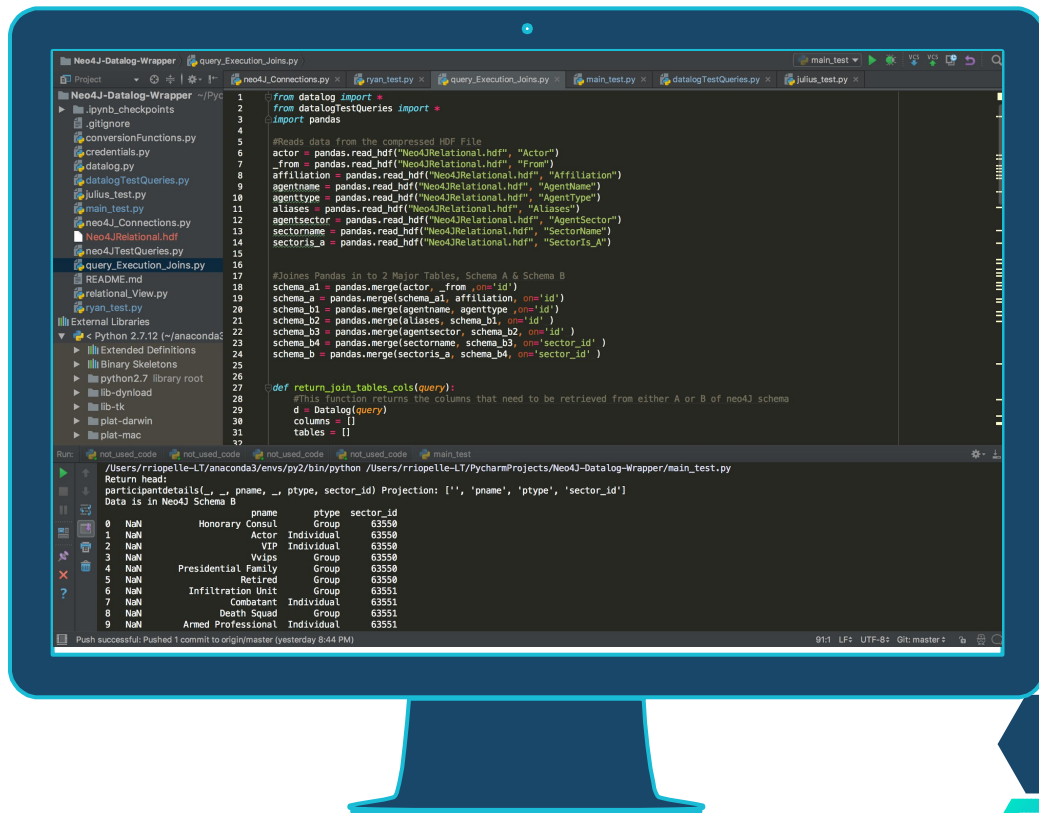


## NEED TO USE GITHUB!

Show and explain your web, app or software projects using these gadget templates.

## PyCharm Really Helps!

Show and explain your web, app or software projects using these gadget templates.



A cluster of various hexagonal icons in shades of blue and teal. The icons include a lightbulb, a thumbs up, a network node, a smartphone, a magnifying glass, a gear, and a speech bubble. The central hexagon is the largest and contains the number 4.

4

# Execution Plans

Need for virtualized architecture


# What is a wrapper?

---

- Wrappers are components of DI systems that communicate with the data sources
  - sending queries from higher levels in the system to the sources
  - converting replies to a format that can be manipulated by query processor
- Complexity of wrapper depends on nature of data source
  - e.g., source is RDBMS, wrapper's task is to interact with JDBC driver
  - in many cases, wrapper must parse semi-structured data such as HTML pages and transform it into a set of tuples
  - we focus on this latter case

# Datalog Parser Overview

- Class Datalog (datalog.py)
- Input: Standard Datalog
- Strategy: regular expressions
- Output: Datalog Object






# Datalog Classes

---

Parsers consists of a set of classes and a super class

- DatalogBase(*object*)
- DatalogAttribute(*DatalogBase*)
- DatalogRelation(*DatalogBase*)
- Datalog(*Base*)

# Datalog NameSpace and Members



# Inspecting Datalog Object Schema

```
query: A(a, b , e) :- mytable(a, b, c, d,
    _), other(f, g, h), CONTAINS(g,'foo'),a
    > 1, GROUP_BY([a, b], d = COUNT(c)),
    d < 100, SORT_BY(b, 'DESC'), f =
    FUN(e), LIMIT(25), DISTINCT
```

## **def inspect(query):**

# use to inspect object properties

d = Datalog(query)

print 'query:', query

print 'head:',d.head

print 'name:',d.name

print 'projection:', d.getList

print 'Join keys:',d.joinKeys

for x in d.relations:

print x.name

print 'predicate:', x.predicateToList

for a in x.attributes:

print '\t', x.attributes[a].name, '\t', \

x.attributes[a].index, '\t', \

x.attributes[a].isJoinPart, '\t', \

x.predicate[x.attributes[a].index] if x.attributes[a].index in x.predicate else

''

print 'group by:', d.groupBy

print 'grouping columns:', d.groupByColumns

print 'aggregations:', d.aggregations

print 'having clause:', d.predicateToList

print 'order by:', d.orderBy

print 'limit:', d.limit

print 'distinct:', d.distinct

# Sample Output

head: A(a, b , e)

name: A

projection: ['a', 'b', 'e']

Join keys: []

mytable

predicate: []

a 0 False a > 1

c 2 False

b 1 False

d 3 False d < 100

other

predicate: ["g = CONTAINS('foo')"]

h 2 False

g 1 False CONTAINS('foo')

f 0 False f = FUN(e)

group by: ['GROUP\_BY([a, b], d =  
COUNT(c))']

grouping columns: ['a', 'b']

aggregations: ['COUNT(c) AS d']

having clause: []

order by: ORDER BY b DESC

limit: 25


distinct: True

# Solution

---

- Standardized syntax makes things easier
  - All functions represented in CAPS
  - All relations represented in lowercase
- Limiting use of nesting to a few cases
  - Write special code to identify these cases and handle appropriately

# Connection Strings



# Accessing Neo4J Using Python

- Neo4j Python Driver
- The Example Project
- Neo4j Community Drivers
- Py2neo
- Neo4jRestClient
- Bulbflow

```
from neo4j.v1 import GraphDatabase, basic_auth
driver = GraphDatabase.driver("bolt://54.85.112.231:7687",
auth=basic_auth("neo4j", "LEbKqX3q"))
session = driver.session()
```

```
authenticate("54.85.112.231:7474", "neo4j", "LEbKqX3q")
graph = Graph("bolt://54.85.112.231/db/data/")
```

# Query Execution Classes

---

Parsers consists of a set of classes and a super class

- ReturnJoinTablesCols (*query*)
- ReturnSchemaA\_or\_B(*tables, columns*)
- ProjectedDataOutput(*dataframe*)
- ExecuteQuery(query)



# Cypher Converted to Relational View

## Cymer Relational View

- Actor: [u'AliasList' u'id' u'pname' u'ptype']
- From: [u'country' u'id']
- Affiliation: [u'end' u'id' u'org' u'start']
- AgentName: [u'id' u'pname']
- AgentType: [u'id' u'ptype']
- Aliases: [u'alias' u'id']
- AgentSector: [u'id' u'sector\_id']
- SectorName: [u'name' u'sector\_id']
- SectorIs\_A: [u'sector\_id' u'sector\_id2']

## Compared to MySQL

- Single GTD Table

## Compared to Postgre

- Two GDELT Tables

# Ideal State w/ ID

*DLOG*:  $q(\text{name}) :- \text{actor}(\text{id}, \_, \text{name}, \_), \text{affiliation}(\text{id}, \text{'Taliban'}, \_, \_)$

*Map*:  $\langle \text{relation} \rangle \rightarrow \langle \text{node/edge} \rangle$

$\text{actor}(\text{id}, \text{name}) \rightarrow \text{ID}(\langle \text{n} \rangle \text{Actor}), \langle \text{n} \rangle \text{Actor.Name},$

$\text{affiliation}(\text{id}, \text{'Taliban'}) \rightarrow \text{ID}(\langle \text{n} \rangle \text{Actor}), \langle \text{n} \rangle \text{Organization.Name}$

*Cypher*:

`MATCH (o: Organization {Name: 'Taliban'})-[]-(p:Actor)`

`RETURN p.Name as name`

# Ideal State Arbitrary Join

*DLOG:*  $q(\text{name}) \text{ :- actor}(\_, \_, \text{name}, \_), \text{agentname}(\_, \text{pname})$

*Map:*  $\langle \text{relation} \rangle \rightarrow \langle \text{node/edge} \rangle$

$\text{actor}(\text{id}, \text{name}) \rightarrow \text{ID}(\langle \text{n} \rangle \text{Actor}), \langle \text{n} \rangle \text{Actor.Name},$

$\text{agentname}(\text{id}, \text{'Taliban'}) \rightarrow \text{ID}(\langle \text{n} \rangle \text{AgentName}), \langle \text{n} \rangle \text{Organization.Name}$

*Cypher:*

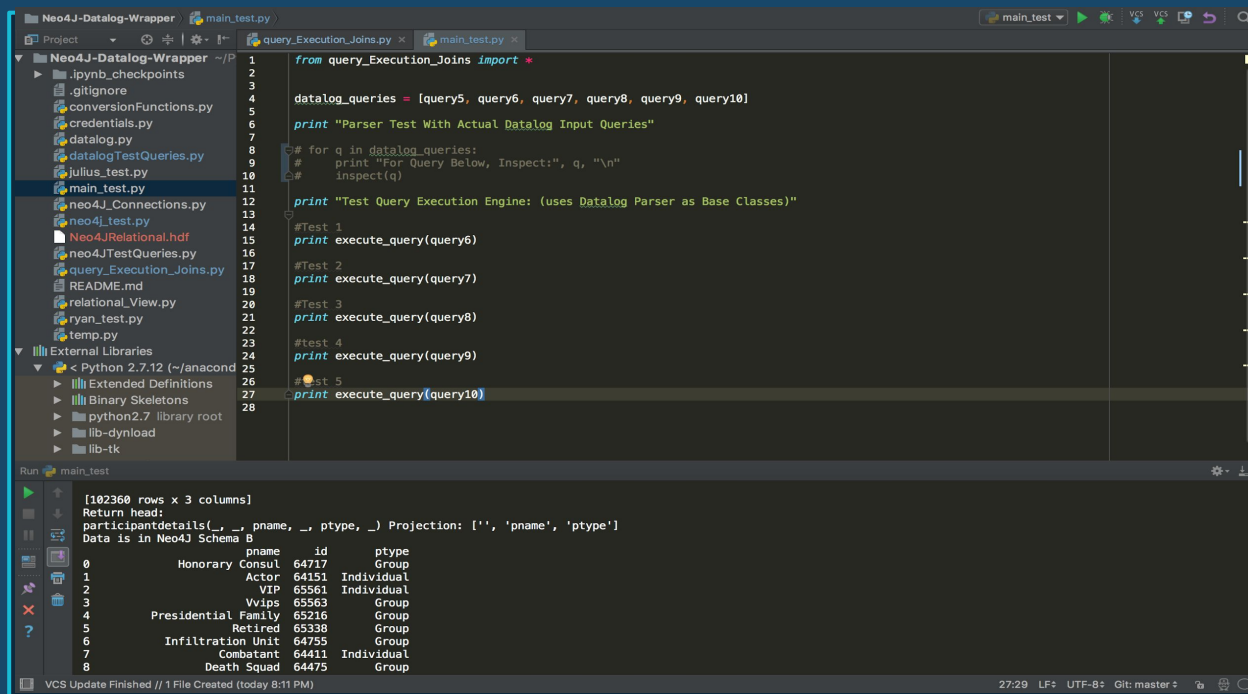
MATCH (o: AgentName)

MATCH (p: Actor) WHERE o.Name = p.Name

RETURN p.Name as name

## Example Code

Show and explain your web, app or software projects using these gadget templates.



The screenshot shows a Python IDE with a project named 'Neo4J-Datalog-Wrapper'. The project structure includes files like 'main\_test.py', 'neo4j\_Connections.py', 'neo4j\_test.py', 'Neo4JRelational.hdf', 'neo4jTestQueries.py', 'query\_Execution\_Joins.py', 'README.md', 'relational\_View.py', 'ryan\_test.py', and 'temp.py'. The 'main\_test.py' file is open, showing a script that imports 'query\_Execution\_Joins' and defines a list of datalog queries. The script then tests the execution of these queries using the 'execute\_query' function. The output of the script is displayed in the 'Run' console, showing a table of results with columns 'pname', 'id', and 'ptype'.

```
from query_Execution_Joins import *

datalog_queries = [query5, query6, query7, query8, query9, query10]

print "Parser Test With Actual Datalog Input Queries"

# for q in datalog_queries:
#     print "For Query Below, Inspect:", q, "\n"
#     inspect(q)

print "Test Query Execution Engine: (uses Datalog Parser as Base Classes)"

#Test 1
print execute_query(query6)

#Test 2
print execute_query(query7)

#Test 3
print execute_query(query8)

#Test 4
print execute_query(query9)

#Test 5
print execute_query(query10)
```

Run main\_test

```
[102360 rows x 3 columns]
Return head:
participantdetails(_ , _ , pname, _ , ptype, _) Projection: ['', 'pname', 'ptype']
Data is in Neo4J Schema B
```

	pname	id	ptype
0	Honorary Consul	64717	Group
1	Actor	64151	Individual
2	VIP	65561	Individual
3	Vvips	65563	Group
4	Presidential Family	65216	Group
5	Retired	65338	Group
6	Infiltration Unit	64755	Group
7	Combatant	64411	Individual
8	Death Squad	64475	Group

VCS Update Finished // 1 File Created (today 8:11 PM) 27:29 LF: UTF-8: Git: master



# Issues/Improvements

Need for virtualized architecture

# Issues

---

Nested expressions in datalog

- difficult to parse using regex

Identifying Functions vs Relations

- Syntactically both look the same
- Not sure if data should have been modeled in two separate databases

# Issue - Example

Group By Example:

```
Pattern = `(\w+([().+[]]))`
```

```
Datalog = `a(x, y, a, b) :- b(x, y, z),  
                           GROUP_BY([x, y], a = COUNT(z), b = SUM(Z))`
```

Some Possible Matches in body:

```
> b(x, y, z)
```

```
> b(x, y, z), GROUP_BY([x, y], a = COUNT(z), b = SUM(Z))`
```

```
> GROUP_BY([x, y], a = COUNT(z)
```

```
> COUNT(z)
```

```
> GROUP_BY([x, y], a = COUNT(z), b = SUM(Z))
```

# Issues

---

- i) Data should not be modeled in two separate schemas
- ii) Cannot use surrogate key IDs for joining or tracking of data
- iii) Data variables in global schema do not equal those in Neo4j or Return statements. I.e. Proposed global projection variables does not make sense (not tied directly to our portion, but didn't seem correct)
- iv) Match statements take too long to run or don't run at all
  - (1) Prevents correct predicate pushdown
  - (2) Provides a reason to use HDF files
- v) Actually creating one single joined Neo4J does not work





# Lessons/Future Considerations

Need for virtualized architecture


# DI Lessons Learned

---


- Regex is good for only the simple cases
  - Lexical parsers like PyParsing handle nesting better and allow for user defined grammar
  - Understanding the graph database schema is necessary to translate
- It's not easy!

# Github Stats

Contributors	Traffic	Commits	Code frequency	Punch card	Network
--------------	---------	---------	----------------	------------	---------




## Git clones



**8**  
Clones

**4**  
Unique cloners

## Visitors



**249**  
Views

**10**  
Unique visitors



# Thanks!

**Any questions?**

